

The Central Distribution: A Minimax-KL Approach to Noninformative Priors and Feature Selection

Behrouz Ayati* 

Abstract. We introduce the *central distribution*, a probability distribution that represents a set, or more generally a hierarchy, of probability distributions by minimizing the worst-case Kullback–Leibler (KL) divergence to the admissible models. This provides a principled method for converting partial knowledge and set-valued uncertainty into a single probability measure suitable for Bayesian inference.

We show that central distributions admit a minimax characterization and can be obtained as solutions to convex–concave saddle point problems. For parametric models, we establish a sufficient fixed-point condition for priors whose induced marginal distributions are central, revealing a close connection with Bernardo’s reference prior methodology. We extend the framework to hierarchical uncertainty structures and propose an asymptotic notion of centrality for infinite sequences.

Finally, we show how uncertainty about relevant feature subsets can be represented hierarchically, and how the resulting central distributions naturally induce a preference for sparse and structurally simpler models. This suggests a unified information-theoretic framework for constructing noninformative priors and performing Bayesian feature selection.

MSC2020 subject classifications: Primary 62F15; secondary 62B10, 62H30.

Keywords: Central distribution, noninformative prior, minimax-KL, feature selection, partial knowledge, hierarchical uncertainty.

1 Introduction

We begin by revisiting a fundamental decision-making problem and its treatment within Bayesian decision theory.

Consider an environment where an agent must select an action from a set of available actions. The loss associated with each action depends on the unknown state of the environment and can be quantified by a real number. Specifically, let a denote an action, and x represent the unknown environment state. The loss function $\ell(a, x)$ quantifies the loss for each action.

We assume that some evidence e about the unknown environment state is available, and the agent can leverage this evidence to determine the best action. In other words, our agent aims to find an optimal strategy, where a strategy defines the distribution of each action conditioned on the observed evidence: $S(A | E)$.

*Independent Researcher, aybehrouz@gmail.com

According to Bayesian decision theory, the optimal strategy is the one that minimizes the expected loss:

$$S_{\text{opt}} = \arg \min_S \sum_a \sum_{x,e} \ell(a, x) S(a | e) P^*(x, e),$$

where P^* is the joint probability distribution of the environment state and the evidence.

One can show that there always exists a deterministic optimal strategy, namely a strategy that, for each evidence e , selects the action minimizing the expected loss conditioned on e :

$$S_{\text{opt}}(e) = \arg \min_a \mathbb{E}_{P^*} [\ell(a, x) | e]. \quad (1)$$

Now, suppose that we decide to use a different evidence E' instead of E . Could this change potentially improve our strategy? Within the framework of Bayesian decision theory, we can easily answer this question: To compare any two strategies, we can simply compare their expected loss. The strategy with the lower expected loss is the better one.

In practice, we cannot choose arbitrary evidence. We can only use evidence that is a deterministic function of the evidence provided by the environment, allowing its value to be determined after observing the original evidence.

Using a simpler form of evidence resembles feature selection in machine learning. For example, if the evidence is an image, we might replace it with the first three Fourier coefficients of the image. Since these coefficients are a function of the original evidence, they satisfy $E' = f(E)$. Under our Bayesian framework, however, the optimal strategy based on E is always at least as good as any strategy based on E' .

One might argue that adopting a strategy based on simpler evidence may discard useful information and is therefore not justified by Bayesian decision theory. The key point is that the space of strategies using E' is a subset of the space of strategies using E . When we are searching for the best strategy based on E , we are implicitly evaluating all simpler forms as well. This helps explain why standard Bayesian inference does not inherently favor simpler evidence representations.

With an appropriate choice of prior distribution, the Bayesian framework can automatically extract the relevant information from the evidence when determining an optimal strategy. In this sense, an appropriate prior could induce feature selection. However, when applied in a conventional Bayesian setting, many commonly used noninformative priors, such as the flat prior, do not naturally promote feature selection.

In this paper, we introduce the central distribution, a minimax-KL representative of a set, or more generally a hierarchy, of probability distributions. By establishing a fixed-point condition for priors whose induced marginals are central, this framework offers a new information-theoretic generalization of objective Bayesian reference analysis [2]. It moves beyond maximizing expected missing information to provide a unified mechanism for minimizing worst-case divergence under layered, set-valued uncertainty. Furthermore, we propose asymptotic characterizations of central priors and investigate the role of hierarchical central distributions in Bayesian feature selection.

1.1 Related Work

The problem of constructing objective or noninformative priors has a long history in Bayesian statistics. Among the most influential approaches are Jeffreys priors, reference priors, and maximum-entropy based methods [5, 1, 4].

Jeffreys introduced a general rule for constructing invariant noninformative priors based on the Fisher information matrix [5]. The resulting priors are widely used because of their parameterization invariance and their connection to asymptotic likelihood theory. However, Jeffreys priors are derived from local geometric properties of the parameter space and do not directly address the problem of representing a set of admissible probability distributions by a single distribution.

The closest connection to the present work is Bernardo’s reference prior methodology [2, 1]. Reference priors are constructed by maximizing the expected information gained from the data and are intended to represent a state of minimal prior information. In Section 2.3, we show that central priors satisfy a fixed-point condition closely related to the reference prior construction. In particular, for a single parameter and a single set of distributions, the resulting central prior coincides with the corresponding reference prior.

Several authors have also studied objective Bayesian inference through information-theoretic principles [4, 6]. Examples include maximum entropy methods, minimum discrimination information principles, and related divergence-based approaches. These methods typically select a distribution by optimizing an information criterion subject to constraints. In contrast, the central distribution is defined as a minimax representative of a set or hierarchy of distributions, minimizing the worst-case Kullback–Leibler divergence to the admissible models.

The present work is also related to robust Bayesian statistics and imprecise probability [7, 8]. In these frameworks uncertainty is represented by sets of probability distributions rather than a single model. Existing approaches generally focus on performing inference directly with sets of distributions. The central distribution framework instead seeks a single representative distribution that summarizes the available set-valued or hierarchical uncertainty while preserving its information-theoretic structure.

To the best of our knowledge, existing objective-prior methodologies do not provide a unified mechanism for representing hierarchical uncertainty structures and converting them into a single probability distribution through a minimax-KL criterion. The central distribution framework is proposed as a step in that direction.

2 Central Distribution

Consider a random variable X . In Bayesian modeling, our knowledge about X is represented by a probability distribution, which serves as a basis for making inferences.

However, in many situations, our information may be too limited to justify choosing a single distribution. Instead, we may be able to specify only a set of plausible distributions. Such a set represents some form of *uncertainty* or *partial knowledge*: although it

does not identify a single distribution, it constrains the range of plausible probabilistic models.

A well-known example of this idea is the exchangeability assumption used in de Finetti’s theorem. A sequence of random variables $X_1, X_2, \dots$ is said to be exchangeable if the ordering of the variables carries no information. That is, the joint probability distribution remains invariant under permutations of the variables. This assumption imposes a constraint, effectively introducing a set of plausible distributions. Thus, exchangeability can be seen as a specific form of partial knowledge.

In this work, we adopt the view that a set of probability distributions can represent a meaningful state of partial knowledge, and that such set-valued uncertainty may itself be used as the basis for inference.

More generally, partial knowledge may possess a hierarchical structure. In the simplest case, uncertainty is represented by a single set of candidate distributions. In richer settings, however, this uncertainty may be organized into nested collections, where elements of a set may themselves be sets of distributions. Such nesting can continue recursively, yielding a hierarchical structure.

Crucially, the choice of hierarchical organization is itself informative: grouping distributions into subsets expresses that certain alternatives are regarded as belonging to a common epistemic category and should therefore be aggregated before being compared with alternatives in other groups. Therefore, the hierarchy encodes information beyond the collection of leaf distributions alone.

This hierarchical organization may be interpreted as introducing multiple layers of latent structural decisions, where higher-level choices restrict the admissible lower-level alternatives.

Within a Bayesian framework, inference ultimately requires a single probability measure. Hence, to use hierarchical partial knowledge in Bayesian inference, we require a principled method for transforming a nested set-valued uncertainty into a single probability distribution while preserving the informational structure encoded by the hierarchy.

We adopt the principle that the probability measure used to represent prior beliefs about the environment should depend only on the available information about the environment itself, and not on agent-specific considerations such as the loss function or the informativeness of observations. We call this modeling assumption the *principle of objectivity*.

From a minimax perspective, when uncertainty is represented by a set of candidate distributions, a natural representative distribution is one that minimizes the worst-case divergence to the members of the set. Such a distribution may be viewed as the most conservative or robust single-distribution summary of the available partial knowledge.

Notation 2.1 (Kullback–Leibler Divergence). For probability distributions P and Q , $\mathbb{D}_{\text{KL}}[Q(X) \parallel P(X)]$ denotes the Kullback–Leibler divergence between P and Q over a random variable X . For a discrete random variable X we have

$$\mathbb{D}_{\text{KL}}[Q(X) \parallel P(X)] = \sum_x Q(x) \ln \frac{Q(x)}{P(x)}.$$

When the underlying random variable is clear from context, we omit the argument and write $\mathbb{D}_{\text{KL}}[Q \parallel P]$.

Definition 2.2. Let X be a random variable and $\mathcal{Q}$ be a set of probability distributions over X . We call a distribution C a *central distribution* of $\mathcal{Q}$ for X if, for every probability distribution P , we have

$$\sup_{Q \in \mathcal{Q}} \mathbb{D}_{\text{KL}}[Q(X) \parallel C(X)] \leq \sup_{Q \in \mathcal{Q}} \mathbb{D}_{\text{KL}}[Q(X) \parallel P(X)], \quad (2)$$

By convention, any probability distribution is regarded as its own central distribution.

Under this definition, the central distribution of $\mathcal{Q}$ is not necessarily an element of $\mathcal{Q}$. That is, the distribution that best represents the set in the specified sense may lie outside the set it summarizes.

Sometimes partial knowledge is represented by a hierarchical structure of sets of distributions. Our current definition of the central distribution does not account for this case. The following recursive definitions extend the concept to handle such situations. While not directly about feature selection, this extension plays a key role in enabling the connection between central distributions and feature selection, a connection that will become clearer later.

Definition 2.3 (Hierarchy of Distributions). A *hierarchy of distributions* over X is defined recursively as follows:

- Any probability distribution over X is a hierarchy of distributions over X .
- Any set whose elements are hierarchies of distributions over X is itself a hierarchy of distributions over X .

Intuitively, a central distribution acts as a representative or summary of a set of distributions. While it may not belong to the set itself, it reflects the overall properties of the set in a way that is meaningful for inference. It also provides a way to evaluate probability distributions under partial knowledge, motivating the definition of a corresponding central divergence.

The concept of a central distribution is not intended to yield a unique object; rather, all distributions satisfying the definition can be regarded as equally good representatives. However, when multiple central distributions exist, our constructions should not depend on an arbitrary choice among them. To avoid such dependence, we define the central divergence using the best divergence attained over the set of all central distributions.

Definition 2.4 (Central Divergence). For any distribution P , we define the *central divergence* of P from a hierarchy of distributions $\mathcal{H}$, with respect to a random variable X , as follows:

$$\mathbb{D}_{\text{C}}[\mathcal{H}(X) \parallel P(X)] = \inf_C \mathbb{D}_{\text{KL}}[C(X) \parallel P(X)],$$

where C denotes a central distribution of $\mathcal{H}$ for X .

So far, this definition applies only when the hierarchy is a single set of distributions. The following definition extends the construction recursively to general hierarchical uncertainty structures.

Definition 2.5 (Central Distribution). Let X be a random variable, and let $\mathcal{H}^*$ be a hierarchy of distributions over X . We call a distribution C a *central distribution* of $\mathcal{H}^*$ if, for every probability distribution P , we have:

$$\sup_{\mathcal{H} \in \mathcal{H}^*} \mathbb{D}_C [\mathcal{H}(X) \parallel C(X)] \leq \sup_{\mathcal{H} \in \mathcal{H}^*} \mathbb{D}_C [\mathcal{H}(X) \parallel P(X)].$$

Definition 2.5 generalizes Definition 2.2, which is recovered when the hierarchy consists of a set of probability distributions.

For later use, we introduce a compact notation for the minimax objective appearing in Definition 2.2.

Notation 2.6. For a set of probability distributions $\mathcal{Q}$ over X , we define the functional

$$\mathbb{S}_{\mathcal{Q}} [P(X)] := \sup_{Q \in \mathcal{Q}} \mathbb{D}_{\text{KL}} [Q(X) \parallel P(X)].$$

When $\mathcal{Q}$ and X are clear from the context, we may simply write $\mathbb{S} [P]$.

This notation will be used throughout the sequel and allows Definition 2.2 to be written compactly as

$$C \in \arg \min_P \mathbb{S}_{\mathcal{Q}} [P(X)].$$

A simple example illustrates that hierarchical organization can affect the resulting representative distribution even when the underlying admissible sets remain unchanged.

Let X be a random variable with three possible outcomes x_1, x_2, x_3 . For each i , let S_i denote the set of all probability distributions over X that assign zero probability to x_i .

By symmetry, the central distribution of each set S_i is the uniform distribution over its support:

$$C_{S_1} = (0, 0.5, 0.5), \quad C_{S_2} = (0.5, 0, 0.5), \quad C_{S_3} = (0.5, 0.5, 0).$$

If the three sets S_1, S_2, S_3 are treated symmetrically as elements of a single set, then their central distribution is

$$C_{\{S_1, S_2, S_3\}} = \left(\frac{1}{3}, \frac{1}{3}, \frac{1}{3} \right).$$

Now consider instead a hierarchical organization in which S_1 and S_2 are grouped together, while S_3 forms its own separate branch. By numerically solving the minimax problem in Definition 2.5, we obtain

$$C_{\{S_1, S_2\}} = (0.25, 0.25, 0.5).$$

Applying the same procedure to the two distributions $C_{\{S_1, S_2\}}$ and C_{S_3} yields

$$C_{\{\{S_1, S_2\}, S_3\}} = (0.4, 0.4, 0.2).$$

Thus, although the underlying admissible sets are identical in both constructions, the resulting representative distribution differs.

This demonstrates that hierarchical structure carries informational content beyond the collection of leaf distributions alone. Consequently, flattening a hierarchy into a single unstructured set may discard meaningful epistemic structure.

In many learning problems, uncertainty concerns not only the parameters of a predictive model but also its structural form. In particular, when the relevance of potential features is unknown, each candidate feature subset determines a distinct set of plausible models in which the target variable depends only on that feature subset.

Feature selection may therefore be viewed as uncertainty over a collection of distributions. The hierarchical central distribution framework provides a principled mechanism for representing and aggregating such uncertainty. By organizing candidate models hierarchically, one can encode preferences over structural hypotheses, such as favoring sparse subsets, grouping related attributes, or privileging particular interactions.

Candidate feature subsets typically induce overlapping families of admissible distributions. Distributions representing simpler dependency structure may belong simultaneously to many such families, whereas more complex dependency patterns belong to fewer. As a result, hierarchical aggregation through central distributions can induce a preference toward structurally simpler models by assigning more weight to distributions compatible with multiple structural hypotheses.

2.1 Minimax Characterization

With the concept of hierarchical central distributions in place, we now develop the theoretical machinery needed to characterize and compute them.

Throughout this subsection, all random variables are assumed to be discrete with finite support. Before proceeding, we establish the existence of central distributions in the finite setting considered throughout this section.

Theorem 2.7. *Let X be a discrete random variable with finite support, and let $\mathcal{Q}$ be a finite set of probability distributions over X . Then a central distribution of $\mathcal{Q}$ for X exists.*

Proof. For a fixed distribution Q , we may write

$$\mathbb{D}_{\text{KL}}[Q(X) \parallel P(X)] = c - \sum_x Q(x) \ln P(x).$$

Since $\ln$ is upper semicontinuous on $[0, +\infty)$ and X has finite support, it follows that $\mathbb{D}_{\text{KL}}[Q \parallel P]$ is lower semicontinuous as a function of P .

Because $\mathcal{Q}$ is finite, $\mathbb{S}_{\mathcal{Q}}[P(X)]$ is the maximum of finitely many lower semicontinuous functions and is therefore itself lower semicontinuous.

The set of probability distributions over X is a compact simplex. Hence, by the extreme value theorem, $\mathbb{S}_{\mathcal{Q}}[P]$ attains its minimum. Therefore there exists a probability distribution C such that

$$\mathbb{S}_{\mathcal{Q}}[C(X)] = \inf_P \mathbb{S}_{\mathcal{Q}}[P(X)].$$

Thus C is a central distribution of $\mathcal{Q}$ for X . $\square$

The central distribution is essentially the solution to a convex–concave game. Because solutions of convex–concave games correspond to saddle points, by defining the appropriate objective function, we can characterize the central distribution using standard convex optimization techniques.

Definition 2.8 (Saddle Point). Consider a function $f : \mathcal{W} \times \mathcal{Z} \rightarrow \mathbb{R}$. We refer to a pair (w^*, z^*) as a *saddle point* for f , if we have

$$f(w^*, z) \leq f(w^*, z^*) \leq f(w, z^*),$$

for all $w \in \mathcal{W}$ and $z \in \mathcal{Z}$.

Next, we establish the connection between saddle points and central distributions.

Theorem 2.9. Let X be a discrete random variable, and $\mathcal{Q}$ be a finite set of distributions over X . Consider the following function:

$$f(P, R) = \sum_i \mathbb{D}_{\text{KL}}[Q_i(X) \parallel P(X)] R(i),$$

where P and R are probability distributions, $Q_i \in \mathcal{Q}$ and the index i ranges over all elements of $\mathcal{Q}$.

If (P^*, R^*) is a saddle point for f , then P^* is a central distribution of $\mathcal{Q}$ for X .

Proof. Let C be a central distribution of $\mathcal{Q}$ for X , and (P^*, R^*) be a saddle point for f . Since X is a discrete random variable and $\mathcal{Q}$ is finite, C exists. As C is a probability distribution, the point (C, R^*) lies within the domain of f . By the definition of a saddle point, this implies that

$$f(P^*, R^*) \leq f(C, R^*). \quad (3)$$

Because R^* is a probability distribution, the definition of f implies:

$$f(C, R^*) \leq \mathbb{S}[C].$$

On the other hand, by the definition of a saddle point

$$f(P^*, R^*) = \max_R f(P^*, R),$$

where R is a probability distribution, and therefore

$$\max_R f(P^*, R) = \mathbb{S}[P^*].$$

Thus, Equation 3 implies

$$\mathbb{S}[P^*] \leq \mathbb{S}[C].$$

Since C is a central distribution, we must have

$$\mathbb{S}[P^*] = \mathbb{S}[C],$$

and P^* is also a central distribution of $\mathcal{Q}$ for X . $\square$

Lemma 2.10. *Let f be the function defined in Theorem 2.9. The point (P^*, R^*) is a saddle point for f , if and only if there are constants λ , η and a function μ , such that for every i , x the following conditions hold:*

$$\sum_x Q_i(x) \ln \frac{Q_i(x)}{P^*(x)} + \mu(i) = \lambda, \quad (4)$$

$$P^*(x) = \sum_i R^*(i) Q_i(x) \quad (5)$$

$$\sum_i R^*(i) = 1, \quad (6)$$

$$R^*(i) \geq 0, \quad (7)$$

$$\mu(i) \geq 0,$$

$$\mu(i) R^*(i) = 0. \quad (8)$$

Furthermore, for a fixed probability distribution R , a distribution $\hat{P}$ minimizes $f(\cdot, R)$, if and only if for all x it satisfies

$$\hat{P}(x) = \sum_i Q_i(x) R(i). \quad (9)$$

Proof. Similar to f , we define $\hat{f}$ as follows:

$$\hat{f}(P, R) = \sum_i R(i) \sum_x Q_i(x) \ln \frac{Q_i(x)}{P(x)},$$

where the constraints that P and R be probability distributions are omitted. This enables us to formulate the saddle point of f as the solution to a constrained optimization problem defined for $\hat{f}$.

Let (P^*, R^*) be a simultaneous solution to the following two optimization problems:

$$\begin{aligned}
& \text{minimize} && \hat{f}(P, R) && \text{w.r.t. } P \\
& \text{subject to} && \sum_x P(x) = 1, \\
& && && \\
& \text{maximize} && \hat{f}(P, R) && \text{w.r.t. } R \\
& \text{subject to} && R(i) \geq 0, && \text{for all } i \\
& && \sum_i R(i) = 1
\end{aligned} \tag{10}$$

Clearly, if (P^*, R^*) exists, it will be a saddle point for f .

Both of these optimization problems are convex and differentiable in the sense of convex optimization; see [3]. To establish this, consider that all the constraint functions are affine, and therefore are convex, concave and differentiable.

Additionally, $\hat{f}$ is a differentiable function with respect to both R and P . For any fixed P , $\hat{f}(P, \cdot)$ is an affine function and is concave.

On the other hand, for a fixed solution to the second optimization problem denoted by $\hat{R}$, we can write

$$\hat{f}(P, \hat{R}) = c - \sum_{i,x} \hat{R}(i) Q_i(x) \ln P(x), \tag{11}$$

where c is not a function of P . Since $\hat{R}$ is a solution to the second optimization problem, it satisfies the corresponding constraints and is therefore a valid probability distribution. Hence, for all i, x we have $\hat{R}(i) Q_i(x) \geq 0$, and $\hat{f}(\cdot, \hat{R})$ is the negative of a linear combination of concave $\ln P(x)$ functions with nonnegative coefficients. This function is convex.

Thus, both optimization problems are convex and differentiable, with affine constraints. The constraints simply restrict the feasible set to the set of probability distributions, and it is always possible to find feasible points that meet all the constraints. Therefore, Slater's condition is satisfied, and the KKT conditions provide necessary and sufficient conditions for optimality [3]. In other words, if there exists a point (P^*, R^*) that satisfies the KKT conditions for both problems, it will be a saddle point for f , and vice versa.

The KKT conditions are as follows for every x, i and constants λ, η :

$$\sum_i R^*(i) Q_i(x) = \eta P^*(x) \tag{12}$$

$$\sum_x P^*(x) = 1, \tag{13}$$

$$\begin{aligned}
\sum_x Q_i(x) \ln \frac{Q_i(x)}{P^*(x)} + \mu(i) &= \lambda, \\
\sum_i R^*(i) &= 1,
\end{aligned} \tag{14}$$

$$\begin{aligned} R^*(i) &\geq 0, \\ \mu(i) &\geq 0, \\ \mu(i)R^*(i) &= 0. \end{aligned}$$

By summing Equation 12 over x and using Equation 14, we have

$$\eta = 1.$$

Thus, we get Equation 5 which also implicitly implies Equation 13.

To prove the second part of the lemma, we observe that when R is a probability distribution, Equations 12 and 13 are the KKT conditions for the minimization problem in Equation 10. Since this minimization problem is convex and satisfies Slater's condition, the KKT conditions are necessary and sufficient. Hence

$$\hat{P}(x) = \sum_i Q_i(x)R(i)$$

if and only if $\hat{P}$ minimizes $f(\cdot, R)$, where R is any fixed distribution. $\square$

Lemma 2.11. *Let $\mathcal{Q}$ be a finite set of probability distributions over a discrete random variable X . Define the function g as:*

$$g(R) = \sum_i \mathbb{D}_{\text{KL}} [Q_i(X) \parallel R(X)] R(i),$$

where $Q_i \in \mathcal{Q}$, and R is a probability distribution such that

$$R(X, I) = Q_i(X)R(I).$$

Let f be the function defined in Theorem 2.9. If (P^*, R^*) is a saddle point for f , then the joint distribution $Q_i(X)R^*(I)$ is a maximizer of g .

Proof. Let (P^*, R^*) be a saddle point for f . By Lemma 2.10, P^* satisfies Equation 5. Letting $R^*(X, I) = Q_i(X)R^*(I)$ we have

$$g(R^*) = f(P^*, R^*).$$

Now, let R be an arbitrary probability distribution such that $g(R)$ is defined. From the second part of Lemma 2.10, it follows that $g(R)$ indeed corresponds to the minimum of $f(\cdot, R)$, which implies

$$g(R) \leq f(P^*, R).$$

On the other hand, since (P^*, R^*) is a saddle point, we also have

$$f(P^*, R) \leq f(P^*, R^*).$$

Combining the two inequalities, we get:

$$g(R) \leq g(R^*).$$

This shows that $g(R^*)$ is the maximum of g , as desired. $\square$

The following theorem gives a standard sufficient condition for optimality in the KL-minimax setting which is the rationale behind the term “central distribution.” Like the center of a circle, the central distribution is equidistant from the boundary.

Theorem 2.12 (Centrality). *Let X be a discrete random variable, and let $\mathcal{Q}$ be a finite set of distributions over X . Suppose that C is a convex combination of elements of $\mathcal{Q}$, that is,*

$$C(x) = \sum_i Q_i(x) R(i),$$

where R is a probability distribution over the indices of $\mathcal{Q}$. Assume that there exists a constant λ such that

$$\mathbb{D}_{\text{KL}} [Q_i(X) \parallel C(X)] = \lambda$$

whenever $R(i) > 0$, and

$$\mathbb{D}_{\text{KL}} [Q_i(X) \parallel C(X)] \leq \lambda$$

whenever $R(i) = 0$. Then C is the unique central distribution of $\mathcal{Q}$ for X .

Proof. By defining

$$\mu(i) = \lambda - \mathbb{D}_{\text{KL}} [Q_i \parallel C],$$

we have $\mu(i) \geq 0$ and $\mu(i)R(i) = 0$ for every i . Therefore, by Lemma 2.10, the point (C, R) is a saddle point, and Theorem 2.9 implies that C is a central distribution of $\mathcal{Q}$ for X .

To prove uniqueness, let C' be an arbitrary probability distribution. We have

$$\begin{aligned} \mathbb{D}_{\text{KL}} [C \parallel C'] &= \sum_i R(i) (\mathbb{D}_{\text{KL}} [Q_i \parallel C'] - \mathbb{D}_{\text{KL}} [Q_i \parallel C]) \\ &= \sum_i \mathbb{D}_{\text{KL}} [Q_i \parallel C'] R(i) - \lambda \\ &\leq \mathbb{S} [C'] - \lambda. \end{aligned} \tag{15}$$

Suppose that C' is also a central distribution of $\mathcal{Q}$ for X . Since C and C' are both central distributions,

$$\mathbb{S} [C'] = \mathbb{S} [C] = \lambda.$$

Therefore, by non-negativity of the KL divergence, we have

$$\mathbb{D}_{\text{KL}} [C \parallel C'] = 0,$$

which implies C and C' are identical probability distributions. $\square$

Corollary 2.13. *Assume the conditions of Theorem 2.12 and let λ be the corresponding constant. Then, for any probability distribution P ,*

$$\mathbb{D}_{\text{C}} [\mathcal{Q} \parallel P] \leq \mathbb{S} [P] - \lambda.$$

Proof. By Theorem 2.12, C is the unique central distribution of $\mathcal{Q}$, and

$$\mathbb{D}_C[\mathcal{Q} \| P] = \mathbb{D}_{\text{KL}}[C \| P].$$

The result then follows immediately from Equation 15. $\square$

Corollary 2.14. *Under the assumptions of Theorem 2.12, C is a maximizer of the function g defined in Lemma 2.11.*

Proof. In the proof of Theorem 2.12, we showed that (C, R) is a saddle point for f . By letting $C(X, I) = Q_i(X)R(I)$, Lemma 2.11 implies that C is a maximizer of the function g . $\square$

Corollary 2.15. *Assuming the conditions of Theorem 2.12 and using the same λ , let P be a convex combination of elements of $\mathcal{Q}$, i.e.,*

$$P(x) = \sum_i Q_i(x)R(i),$$

where R is a probability distribution. Then there exists an index k in the support of R (that is, $R(k) > 0$) such that

$$\lambda \geq \mathbb{D}_{\text{KL}}[Q_k(X) \| P(X)].$$

Proof. We proceed by contradiction. Suppose that for every k in the support of R , we have $\lambda < \mathbb{D}_{\text{KL}}[Q_k \| P]$. That implies

$$g(P) > \lambda = g(C),$$

where C is the distribution specified in Theorem 2.12, and g is the function defined in Lemma 2.11. Note that $g(P)$ is defined, because P is a convex combination of elements of $\mathcal{Q}$.

However, by Corollary 2.14, the distribution C maximizes g , which contradicts the inequality $g(P) > g(C)$. $\square$

2.2 Approximate Central Distributions

Suppose we want to approximate the central distribution of a hierarchy of distributions $\mathcal{H}$. To do so, we need a measure that quantifies the quality of an approximation. A natural choice is the central divergence introduced in Definition 2.4. Since central distributions are intended to serve as representatives of hierarchies of distributions, central divergence may be viewed as a measure of how well a probability distribution represents a hierarchy.

At first glance, computing this quantity appears to require explicit knowledge of the central distribution of $\mathcal{H}$. However, it turns out that the central divergence admits an upper bound that can be estimated without explicitly knowing the central distribution.

In variational settings, it is often more convenient to work with unnormalized measures rather than normalized distributions. Therefore, we state the next theorem in terms of unnormalized measures. We assume throughout that $\tilde{P}(x) > 0$ whenever $Q(x) > 0$ for some $Q \in \mathcal{Q}$, so that all logarithms and KL divergences are well-defined.

Theorem 2.16. *Let $\mathcal{Q}$ be a finite set of probability distributions over a discrete random variable X , and suppose the assumptions of Theorem 2.12 hold. Let $\tilde{P}$ be a linear combination of elements of $\mathcal{R} \subseteq \mathcal{Q}$ with strictly positive weights, and let P be the corresponding normalized distribution. Then*

$$\mathbb{D}_C[\mathcal{Q} \| P] \leq \sup_{Q \in \mathcal{R}} \mathbb{F}[Q, \tilde{P}] - \inf_{Q \in \mathcal{Q}} \mathbb{F}[Q, \tilde{P}],$$

where

$$\mathbb{F}[Q, \tilde{P}] = \mathbb{H}[Q(X)] + \mathbb{E}_Q[\ln \tilde{P}(X)].$$

Proof. By Theorem 2.12, C is the unique central distribution of $\mathcal{Q}$, and therefore from Corollary 2.13, we have

$$\mathbb{D}_C[\mathcal{Q} \| P] \leq \mathbb{S}_{\mathcal{Q}}[P] - \lambda.$$

On the other hand, since P is a convex combination of elements of $\mathcal{R}$ with strictly positive weights, Corollary 2.15 implies that

$$\lambda \geq \inf_{Q \in \mathcal{R}} \mathbb{D}_{\text{KL}}[Q \| P].$$

Combining these two inequalities yields

$$\mathbb{D}_C[\mathcal{Q} \| P] \leq \mathbb{S}_{\mathcal{Q}}[P] - \inf_{Q \in \mathcal{R}} \mathbb{D}_{\text{KL}}[Q \| P]. \quad (16)$$

Now let $\tilde{P}(X) = ZP(X)$, where $Z > 0$ is a normalization constant. Then, for any Q ,

$$\mathbb{D}_{\text{KL}}[Q \| P] = \ln Z - \mathbb{F}[Q, \tilde{P}].$$

Since $\ln Z$ does not depend on Q , we have

$$\begin{aligned} \inf_{Q \in \mathcal{R}} \mathbb{D}_{\text{KL}}[Q \| P] &= \ln Z - \sup_{Q \in \mathcal{R}} \mathbb{F}[Q, \tilde{P}] \\ \sup_{Q \in \mathcal{Q}} \mathbb{D}_{\text{KL}}[Q \| P] &= \ln Z - \inf_{Q \in \mathcal{Q}} \mathbb{F}[Q, \tilde{P}], \end{aligned}$$

and the result follows from Equation 16. $\square$

Since central distributions are defined through marginal distributions, their construction may be interpreted as combining prior specification and Bayesian prediction into a single optimization problem. As we shall see in a later section, this perspective extends to infinite sequences through suitable limiting constructions. A key question is whether central distributions for finite initial segments S_n converge sufficiently rapidly to the corresponding central distribution of the underlying infinite sequence. Establishing such approximation results would make it possible to perform inference by approximating a tractable central distribution while retaining the objectivity properties associated with the infinite-sequence formulation.

2.3 Parametric Families and Central Priors

The results of the previous subsection were established under the assumption that $\mathcal{Q}$ is finite. Many of the underlying constructions, however, naturally extend to parametric families where $\mathcal{Q}$ can be indexed by a finite-dimensional real vector Θ :

$$\mathcal{Q} = \{Q_\theta \mid \theta \in \Theta\}.$$

We do not pursue the most general extension here. Instead, we use this parametric formulation to motivate the notion of a central prior and to study its relationship with central distributions. We can treat Θ as a random variable and interpret Q_θ as a conditional distribution $Q(x \mid \theta)$. In this setting, the weight function $R(i)$ appearing in the finite case becomes a prior distribution over Θ .

Definition 2.17 (Central Prior). Let X be a discrete random variable, and let $\mathcal{Q}$ be a set of probability distributions over X , indexed by a parameter Θ . A nonnegative measure c on Θ is called a central prior of $\mathcal{Q}$ for X if,

$$C(x) = \int Q(x \mid \theta) c(\theta) d\theta$$

is a central distribution of $\mathcal{Q}$ for X .

Definition 2.18 (Interior Distribution). Let $\mathcal{Q}$ be a set of probability distributions over a random variable X . A distribution R is called an *interior distribution* of $\mathcal{Q}$ iff the sets $\mathcal{Q}$ and $\mathcal{Q} \cup \{R\}$ have the same central distributions for X .

Theorem 2.19. Suppose that the conditions of Theorem 2.12 hold, and let C denote the unique central distribution of $\mathcal{Q}$ (as defined therein). If for some $Q_k \in \mathcal{Q}$

$$\mathbb{D}_{\text{KL}}[Q_k(X) \parallel C(X)] < \lambda,$$

then Q_k is an interior distribution of $\mathcal{Q}$.

Proof. Since $\mathbb{D}_{\text{KL}}[Q_k \parallel C] < \lambda$, we must have $R(k) = 0$. Hence, C is a convex combination of elements of $\mathcal{Q} \setminus \{Q_k\}$ and the conditions of Theorem 2.12 remain satisfied for $\mathcal{Q} \setminus \{Q_k\}$. Thus, C is the unique central distribution of $\mathcal{Q} \setminus \{Q_k\}$ and $\mathcal{Q}$. $\square$

Theorem 2.20. Let X be a random variable and let $\mathcal{Q}$ be a set of probability distributions over X . Any distribution that can be expressed as a convex combination of elements of $\mathcal{Q}$ is an interior distribution of $\mathcal{Q}$.

Proof. Let $R = \sum_i \alpha_i Q_i$, where $Q_i \in \mathcal{Q}$, $\alpha_i \geq 0$, and $\sum_i \alpha_i = 1$.

Fix an arbitrary probability distribution P . By the convexity of KL divergence in its first argument, we have:

$$\mathbb{D}_{\text{KL}}[R \parallel P] \leq \sum_i \alpha_i \mathbb{D}_{\text{KL}}[Q_i \parallel P].$$

Now, note that the right-hand side is a weighted average of the KL divergences. So, there must be at least one index j such that

$$\mathbb{D}_{\text{KL}}[R \parallel P] \leq \mathbb{D}_{\text{KL}}[Q_j \parallel P].$$

Since $Q_j \in \mathcal{Q}$, that implies

$$\mathbb{S}_{\mathcal{Q}}[P] = \mathbb{S}_{\mathcal{Q} \cup \{R\}}[P].$$

Therefore, $\mathcal{Q}$ and $\mathcal{Q} \cup \{R\}$ have the same central distributions. $\square$

One important consequence of this theorem is that, in order to identify a central distribution within a parametric model, it suffices to restrict attention to the set of likelihood distributions induced by different values of the model parameters.

Theorem 2.21. *Let X and T be two random variables, and let $\mathcal{Q}$ be a set of distributions over the pair (X, T) . Assume that for every $Q \in \mathcal{Q}$, we have*

$$Q(X, T) = R(X \mid T)Q(T),$$

where R is a fixed conditional distribution. If C is a central distribution of $\mathcal{Q}$ for T , then $R(X \mid T)C(T)$ is a central distribution of $\mathcal{Q}$ for (X, T) .

Proof. Let C be a central distribution of $\mathcal{Q}$ for T . For any arbitrary probability distribution P , the definition of a central distribution implies

$$\mathbb{S}[R(X \mid T)C(T)] \leq \mathbb{S}[R(X \mid T)P(T)]. \quad (17)$$

Since the KL divergence is nonnegative, for every probability distribution P and every $Q \in \mathcal{Q}$ we have

$$\sum_t \mathbb{D}_{\text{KL}}[R(X \mid t) \parallel P(X \mid t)] Q(t) \geq 0.$$

That implies

$$\mathbb{D}_{\text{KL}}[Q(X, T) \parallel R(X \mid T)P(T)] \leq \mathbb{D}_{\text{KL}}[Q(X, T) \parallel P(X, T)].$$

Taking the supremum over $Q \in \mathcal{Q}$ preserves the inequality, since it holds pointwise for every $Q \in \mathcal{Q}$:

$$\mathbb{S}[R(X \mid T)P(T)] \leq \mathbb{S}[P(X, T)].$$

Combining this inequality with Equation 17, we conclude that $R(X \mid T)C(T)$ is a central distribution of $\mathcal{Q}$ for (X, T) . $\square$

The next theorem highlights a formal connection between the central distribution framework and Bernardo's reference prior methodology [2].

Theorem 2.22. *Let X be a discrete random variable, and let $\mathcal{Q}$ be a set of probability distributions over X , indexed by a finite-dimensional vector Θ . Let c be a strictly positive prior distribution for Θ such that for all $Q(x | \theta) \in \mathcal{Q}$ we have*

$$c(\theta) \propto \exp \left\{ \sum_x Q(x | \theta) \ln c(\theta | x) \right\}. \quad (18)$$

Then, c is a central prior of $\mathcal{Q}$ for X .

Proof. From Equation 18, it follows that there exists a constant λ such that

$$c(\theta) = \exp \left\{ -\lambda + \sum_x Q(x | \theta) \ln \frac{Q(x | \theta)c(\theta)}{C(x)} \right\}.$$

Since $c(\theta)$ is strictly positive,

$$\begin{aligned} c(\theta) &= \exp \{ -\lambda + \mathbb{D}_{\text{KL}} [Q_\theta(X) \| C(X)] + \ln c(\theta) \} \\ &= c(\theta) \exp \{ -\lambda + \mathbb{D}_{\text{KL}} [Q_\theta(X) \| C(X)] \} \end{aligned}$$

implies

$$\mathbb{D}_{\text{KL}} [Q_\theta(X) \| C(X)] = \lambda.$$

This holds for all θ . By using a parametric extension of Theorem 2.12 one can show that C is a central distribution of $\mathcal{Q}$ for X , and by definition, c is a central prior of $\mathcal{Q}$ for X . $\square$

This theorem provides a sufficient fixed-point characterization of central priors; it does not address existence or uniqueness in full generality.

Although this theorem is proved directly from Theorem 2.12, Lemma 2.11 suggests an alternative route to the connection with Bernardo's reference prior methodology. Since reference priors are defined by maximizing missing information [2], the optimization characterization provided by Lemma 2.11 also offers another perspective on this relationship.

2.4 Asymptotic Central Distributions

Many probabilistic models are naturally formulated in terms of a sequence of random variables whose length is not fixed in advance. For example, when a coin is tossed repeatedly, there is no natural upper bound on the number of tosses, and any finite sample can be viewed as an initial segment of an infinite sequence.

Motivated by such examples, let S_n denote the first n variables of an infinite sequence of random variables:

$$S_n = (X_1, X_2, \dots, X_n).$$

The collection $\{S_n\}_{n=1}^\infty$ therefore forms a sequence of random variables.

In many settings, our uncertainty about S_n can be represented by a central distribution. A prominent example arises when the sequence is assumed to be exchangeable. As discussed earlier, exchangeability can be viewed as a form of partial knowledge that defines a set of plausible distributions, from which a central distribution for S_n can be constructed.

However, this central distribution generally depends on the sequence length n . According to the principle of objectivity, our beliefs should depend only on the available information and not on arbitrary aspects of the experimental setup, such as the number of variables under consideration.

Importantly, a distribution defined over a longer sequence induces marginal distributions over all shorter initial segments. This suggests that, rather than constructing a separate central distribution for each S_n , we should seek a single probability measure over the infinite sequence, from which the distributions of all finite initial segments can be obtained as marginals.

This motivates the construction of a central distribution that is well defined over S_∞ .

Definition 2.23. Let $\{S_n\}_{n=1}^\infty$ be a sequence of random variables and $\mathcal{Q}$ be a set of probability measures, where each element of $\mathcal{Q}$ defines a distribution over each S_n . Let C be a probability measure that defines a distribution over every S_n . C is a central distribution of $\mathcal{Q}$ for $\{S_n\}_{n=1}^\infty$, if we have

$$\lim_{n \rightarrow \infty} \mathbb{D}_C [Q(S_n) \parallel C(S_n)] = 0.$$

Intuitively, this definition states that as $n \rightarrow \infty$ the distribution that C induces over S_n becomes increasingly similar to a central distribution for S_n . In other words, C is a central distribution for the infinite sequence S_∞ .

Our next goal is to extend Theorem 2.12 so that it also applies in this asymptotic setting, thereby providing a characterization of central distributions for infinite sequences.

Theorem 2.24 (Asymptotic Centrality). *Let $\{S_n\}_{n=1}^\infty$ be a sequence of random variables and $\mathcal{Q}$ be a set of probability measures, where each element of $\mathcal{Q}$ defines a distribution over each S_n . Assume that there exists a natural number M such that, for every $n \geq M$, there exists a distribution C_n satisfying the conditions of Theorem 2.12 for S_n .*

Let C be a probability measure that defines a distribution over every S_n . If there exists a sequence of constants $\{\lambda_n\}_{n=1}^\infty$ such that

$$\lim_{n \rightarrow \infty} \sup_{Q \in \mathcal{Q}} |\mathbb{D}_{\text{KL}} [Q(S_n) \parallel C(S_n)] - \lambda_n| = 0,$$

then C is a central distribution of $\mathcal{Q}$ for $\{S_n\}_{n=1}^\infty$.

Proof. For $n \geq M$, Theorem 2.12 implies that C_n is the unique central distribution of $\mathcal{Q}$ for S_n . Let μ_n be the corresponding constant from Theorem 2.12. By the definition

of a central distribution, we have

$$\mu_n \leq \mathbb{S}[C(S_n)],$$

and, by Corollary 2.15, it also follows that

$$\mu_n \geq \inf_{Q \in \mathcal{Q}} \mathbb{D}_{\text{KL}}[Q(S_n) \| C(S_n)].$$

On the other hand, clearly

$$\left| \inf_{Q \in \mathcal{Q}} \mathbb{D}_{\text{KL}}[Q(S_n) \| C(S_n)] - \lambda_n \right| \leq \sup_{Q \in \mathcal{Q}} |\mathbb{D}_{\text{KL}}[Q(S_n) \| C(S_n)] - \lambda_n|,$$

and

$$|\mathbb{S}[C(S_n)] - \lambda_n| \leq \sup_{Q \in \mathcal{Q}} |\mathbb{D}_{\text{KL}}[Q(S_n) \| C(S_n)] - \lambda_n|.$$

Hence, both $\{\inf_{Q \in \mathcal{Q}} \mathbb{D}_{\text{KL}}[Q(S_n) \| C(S_n)] - \lambda_n\}_{n=1}^{\infty}$ and $\{\mathbb{S}[C(S_n)] - \lambda_n\}_{n=1}^{\infty}$ converge to 0. By the squeeze theorem, the sequence $\{\mu_n - \lambda_n\}_{n=1}^{\infty}$ also converges to 0. We may write

$$\begin{aligned} \lim_{n \rightarrow \infty} (\mathbb{S}[C(S_n)] - \mu_n) &= \lim_{n \rightarrow \infty} (\mathbb{S}[C(S_n)] - \mu_n + \lambda_n - \lambda_n) \\ &= \lim_{n \rightarrow \infty} (\mathbb{S}[C(S_n)] - \lambda_n) - \lim_{n \rightarrow \infty} (\mu_n - \lambda_n) \\ &= 0. \end{aligned}$$

By Corollary 2.13, we have

$$\mathbb{D}_{\text{C}}[\mathcal{Q}(S_n) \| C(S_n)] \leq \mathbb{S}[C(S_n)] - \mu_n,$$

and the squeeze theorem implies

$$\lim_{n \rightarrow \infty} \mathbb{D}_{\text{C}}[\mathcal{Q}(S_n) \| C(S_n)] = 0.$$

Thus, by Definition 2.23, C is a central distribution of $\mathcal{Q}$ for $\{S_n\}_{n=1}^{\infty}$. $\square$

Theorem 2.24 establishes the conditions for asymptotic centrality; however, proving general existence and uniqueness for infinite sequences remains an open problem. As such, we present the following construction as a conjectured framework. Under standard regularity conditions, this approach motivates the functional form of central priors and serves as a conceptual bridge to stimulate future formal proofs. A similar asymptotic construction also appears in [2].

Conjecture 2.25. *Let $\{T_n\}_{n=1}^{\infty}$ be a sequence of random variables, and let $\mathcal{Q}$ be a set of probability measures, where each element of $\mathcal{Q}$ defines a distribution over each T_n . Assume that $\mathcal{Q}$ can be indexed by a finite-dimensional vector Θ , and $\{T_n\}_{n=1}^{\infty}$ is a consistent sequence of estimators for Θ . Define*

$$c(\theta) \propto \lim_{n \rightarrow \infty} \frac{P(\theta | T_n = t_n)_{|t_n=\theta}}{\mu_n}, \quad (19)$$

where $\{\mu_n\}_{n=1}^\infty$ is a sequence of positive constants and $P(\theta \mid T_n)$ denote the posterior distribution obtained from an arbitrary smooth positive prior. If this limit exists and the convergence is uniform over θ , under the usual regularity conditions, $c(\theta)$ is a central prior of $\mathcal{Q}$ for $\{T_n\}_{n=1}^\infty$.

Let k denote the normalizing constant converting proportionality into equality. Since the convergence is uniform and the logarithm is a continuous function, Equation 19 implies

$$\lim_{n \rightarrow \infty} \sup_{\theta} \left| \ln \frac{P(\theta \mid T_n = t_n)|_{t_n=\theta}}{c(\theta)} - \ln \mu_n + \ln k \right| = 0.$$

Because T_n is a consistent estimator for every θ , the sampling distribution $Q(T_n \mid \theta)$ becomes increasingly concentrated around θ as $n \rightarrow \infty$. Therefore, we may write

$$\lim_{n \rightarrow \infty} \sup_{\theta} \left| \sum_{t_n} Q(t_n \mid \theta) \ln \frac{P(\theta \mid t_n)}{c(\theta)} - \ln \mu_n + \ln k \right| = 0.$$

Furthermore, under the usual regularity conditions, the posterior distribution becomes asymptotically insensitive to the choice of smooth positive prior. Motivated by this asymptotic prior-insensitivity, we heuristically replace $P(\theta \mid T_n)$ with $c(\theta \mid T_n)$, obtaining

$$\lim_{n \rightarrow \infty} \sup_{\theta} \left| \sum_{t_n} Q(t_n \mid \theta) \ln \frac{Q(t_n \mid \theta)}{C(t_n)} - \ln \mu_n + \ln k \right| = 0.$$

Letting $\lambda_n = \ln \mu_n - \ln k$, we observe that

$$\lim_{n \rightarrow \infty} \sup_{\theta} |\mathbb{D}_{\text{KL}}[Q(T_n \mid \theta) \parallel C(T_n)] - \lambda_n| = 0.$$

With appropriate additional assumptions, the conditions of Theorem 2.24 are satisfied. Thus, $c(\theta)$ is expected to act as a central prior for the estimator sequence $\{T_n\}_{n=1}^\infty$.

A central distribution over an estimator sequence cannot by itself be used to perform inference at the level of observed outcomes. To do so, we must convert it into a distribution over the primary variables of interest. Theorems 2.21 provides a simple method for this conversion.

Example (Binomial Experiment). Consider a binomial experiment consisting of n independent Bernoulli trials. Let $\mathcal{Q}$ denote the family of probability distributions corresponding to different values of the success probability $\theta \in [0, 1]$, with θ serving as the indexing parameter for $\mathcal{Q}$. Let T_n be the sample proportion of successes in the n trials. Since T_n is a consistent estimator of θ , we can apply Equation 19 to determine a central distribution for the sequence $\{T_n\}_{n=1}^\infty$.

Assuming a uniform prior on θ , Bayes' rule gives the posterior distribution

$$P(\theta \mid T_n = \frac{k}{n}) = (n+1) \binom{n}{k} \theta^k (1-\theta)^{n-k},$$

which is a Beta($k + 1, n - k + 1$) distribution. Using Stirling's formula, we derive its asymptotic form:

$$P\left(\theta \mid T_n = \frac{k}{n}\right) \sim \frac{(n+1)n^{(n+\frac{1}{2})}}{\sqrt{2\pi}(n-k)^{(n-k+\frac{1}{2})}k^{(k+\frac{1}{2})}} \theta^k (1-\theta)^{n-k}.$$

Evaluating the asymptotic expression at $\frac{k}{n} = \theta$, we obtain

$$P(\theta \mid T_n = t)|_{t=\theta} \sim \frac{n+1}{\sqrt{2\pi n}} \theta^{-\frac{1}{2}} (1-\theta)^{-\frac{1}{2}}.$$

Therefore, if we define

$$\mu_n = \frac{n+1}{\sqrt{2\pi n}},$$

the following limit exists and the convergence is uniform over $\theta \in (0, 1)$:

$$\lim_{n \rightarrow \infty} \frac{P(\theta \mid T_n = t)|_{t=\theta}}{\mu_n} = \theta^{-\frac{1}{2}} (1-\theta)^{-\frac{1}{2}}.$$

Example (Fused Multinomial Model). Let $\{S_n\}_{n=1}^{\infty}$ be a sequence in which each $S_n = (X_1, X_2, \dots, X_n)$ represents n independent trials of an experiment with k possible outcomes, and $\mathcal{Q}$ be a set of probability distributions over S_n indexed by a real vector θ . Let $f : \mathcal{X} \rightarrow \mathcal{Y}$ be a deterministic function that maps each outcome of X_i to a label $Y_i = f(X_i)$, where $\mathcal{Y} = \{1, \dots, m\}$. Assume that for all i and $Q_\theta \in \mathcal{Q}$, if $f(x_i) = f(x'_i)$ we have

$$Q_\theta(x_1, x_2, \dots, x_n) = Q_\theta(x'_1, x'_2, \dots, x'_n).$$

That is, the conditional distribution $Q_\theta(X_i \mid Y_i = y)$ is uniform over the preimage $f^{-1}(y)$.

We can consider $\theta = (\theta_1, \dots, \theta_m)$ to be a probability vector over $\mathcal{Y}$, where each θ_y specifies the probability of observing label y . For a sequence S_n , let n_y denote the number of times label $y \in \mathcal{Y}$ appears. Let $|f^{-1}(y)|$ denote the number of outcomes in $\mathcal{X}$ that are mapped to label y by f . Then the probability of observing S_n under $Q_\theta \in \mathcal{Q}$, denoted $Q(S_n \mid \theta)$, is:

$$Q(S_n \mid \theta) = \prod_{y=1}^m \left(\frac{1}{|f^{-1}(y)|} \cdot \theta_y \right)^{n_y}.$$

Now, consider the estimator sequence $T_n = \frac{1}{n}(n_1, n_2, \dots, n_m)$, which is a consistent estimator for θ . The conditional distribution $Q_\theta(S_n \mid T_n)$ is then given by:

$$R(S_n \mid T_n) = \frac{1}{n!} \prod_{y=1}^m \frac{n_y!}{|f^{-1}(y)|^{n_y}}. \quad (20)$$

Note that this distribution is independent of θ . Therefore, a central distribution for S_n can be obtained from a central distribution for T_n .

As $n \rightarrow \infty$, the distribution of T_n converges to a multivariate Gaussian distribution. Using this fact and Equation 19, we can determine a central prior for $\{T_n\}_{n=1}^\infty$:

$$c(\theta) \propto \prod_{y=1}^m \theta_y^{-\frac{1}{2}}.$$

This is a Dirichlet prior and induces a Dirichlet-multinomial distribution for each T_n :

$$C(T_n) = \frac{\Gamma(\frac{m}{2}) \Gamma(n+1)}{\Gamma(n + \frac{m}{2})} \prod_{y=1}^m \frac{\Gamma(n_y + \frac{1}{2})}{\Gamma(\frac{1}{2}) \Gamma(n_y + 1)}.$$

Note that C is a central distribution for the sequence $\{T_n\}_{n=1}^\infty$, viewed as a probability measure on the infinite sequence T_∞ . The notation $C(T_n)$ denotes the marginal distribution of T_n .

By Theorem 2.21 and Equation 20, we can derive a central distribution for $\{S_n\}_{n=1}^\infty$, which induces the following distribution over each S_n :

$$C(S_n) = \frac{\Gamma(\frac{m}{2})}{\Gamma(n + \frac{m}{2})} \prod_{y=1}^m \frac{\Gamma(n_y + \frac{1}{2})}{\Gamma(\frac{1}{2}) |f^{-1}(y)|^{n_y}}.$$

When dealing with layered ignorance and seeking the central distribution of a hierarchy of distributions, as defined in Definition 2.3, we typically encounter situations in which a given level of the hierarchy contains only finitely many alternatives. In such cases, and under appropriate regularity conditions, if there exists a consistent estimator for the index identifying these alternatives, then the central distribution reduces to an equally weighted average of the corresponding representative central distributions.

3 An Application to Feature Selection

A natural application of hierarchical central distributions arises in feature selection. When the relevant set of features is unknown, each candidate feature set induces a distinct family of admissible distributions. Feature selection can therefore be viewed as uncertainty over a hierarchy of probabilistic models, making it a natural setting for the central distribution framework.

To illustrate this idea, we provide a proof-of-concept experiment using synthetic data. A comprehensive empirical evaluation against established feature-selection methods, such as Spike-and-Slab priors, on real-world datasets lies beyond the scope of this foundational paper and is left for future work.

We consider a classification problem where the goal is to predict a target class, denoted by Y , from a set of observed nominal attributes $(A_1, A_2, \dots, A_n)$. We assume that Y may depend on an unknown subset of the attributes, rather than the full set.

This reflects a form of sparsity in the relationship between the class and the attributes, where irrelevant features are ignored in the true generative model.

We can use the fused multinomial model of Example 2.4 to capture this relationship. As an example, assume that Y depends only on A_2, A_5 . Then, we define the labeling function f as follows:

$$f(A_1, A_2, \dots, A_n, Y) = (A_2, A_5, Y).$$

In a fused multinomial model with this labeling function, given A_2, A_5 , the target class Y is conditionally independent from other attributes.

For every possible subset of attributes, a distinct fused multinomial model of this form is constructed. Each such model corresponds to a set of plausible probability distributions. Collecting these sets yields our overall uncertainty representation, which forms a one-level hierarchy of distributions (Definition 2.3).

Let $C_{\mathcal{H}^*}$ denote a central distribution of the full hierarchy, and let C_{f_i} denote the central distribution of the fused multinomial model associated with the labeling function f_i . As discussed earlier, because the number of candidate models is finite, we have

$$C_{\mathcal{H}^*}(S_n) \propto \sum_i C_{f_i}(S_n).$$

To justify this result, observe that we can construct a consistent estimator for the model indicator I by examining equality relationships among the observed frequencies of outcomes. Specifically, the estimator identifies outcomes that occur with equal frequencies in the sample and returns the indicator corresponding to the model that supports those equalities.


As shown in Equation 1, the optimal Bayesian decision rule depends on the choice of loss function. In this example, we use the zero-one loss, under which the optimal strategy predicts the target class with the highest posterior probability. These posterior probabilities are computed using the central distribution derived in Example 2.4. It turns out that each subset-specific fused multinomial model is assigned its corresponding Jeffreys' prior. Consequently, the resulting predictive distribution can be interpreted as Bayesian model averaging over all candidate feature subsets, with each model equipped with its Jeffreys' prior. The central distribution framework provides a minimax-KL justification for this aggregation.

To see how this approach works in practice, we carried out a synthetic experiment with the following setup. We considered binary attributes and generated each attribute independently with equal probability for 0 and 1. The target class was defined as a deterministic function of a fixed subset of 3 out of 12 total attributes, as shown in Table 1. The same subset of relevant attributes and class assignment function were used across all trials.

For each experiment, we generated a training set of size 20 and a test set of size 60. A larger test set was used to reduce variability in performance estimates, taking advantage of the synthetic nature of the data. To assess robustness to irrelevant features, we progressively removed one attribute at a time from the full set of 12, starting with irrelevant ones and stopping just before removing any of the three relevant attributes.

Table 1: Mappings of 3-bit inputs under class functions Y_1 and Y_2

Input (bits)	Y_1 (Input)	Y_2 (Input)
000	a	A
001	b	B
010	c	A
011	d	C
100	e	D
101	f	D
110	g	D
111	h	E

Figure 1: Accuracy versus number of attributes for 20 training samples. Left: under class assignment function Y_1 . Right: under class assignment function Y_2 .

At each step, we trained and tested two versions of the central distribution: one assuming that the target class is a deterministic function of the attributes (CD-Det), and one allowing the target class to follow a conditional probability distribution given the attributes (CD-Prob). Both versions are constructed by aggregating over 2^{12} fused multinomial models, corresponding to all subsets of the 12 attributes. As a baseline, we also trained a decision tree classifier (DTree). All methods were trained on the same training data and evaluated on the same test data. Each configuration was repeated for 100 independent trials, and we reported the average accuracy (correct predictions divided by total test instances) over these trials.

The results, presented in Figure 1, demonstrate that both central-distribution methods substantially outperform the decision tree on the two classification tasks. Interestingly, the relative difficulty of the tasks differs between the methods. For the central-distribution approach, Y_2 is more challenging than Y_1 , resulting in a larger gap between the deterministic and probabilistic variants. This difficulty arises from strong non-deterministic dependencies between subsets of input bits and the assigned class,




Figure 2: Accuracy versus number of attributes for 10 training samples. Left: under class assignment function Y_1 . Right: under class assignment function Y_2 .

which can create misleading probabilistic associations. In contrast, the decision tree performs worse on Y_1 than on Y_2 . Overall, the deterministic central distribution achieved the highest accuracy, while the probabilistic version produced similar but slightly weaker results.

With 20 training samples, removing irrelevant attributes had little effect on performance. Accuracy remained nearly constant as the number of attributes increased from 3 to 12, and the decision tree showed a similar insensitivity to attribute count. However, with only 10 training samples (Figure 2), accuracy decreased noticeably as additional irrelevant attributes were introduced, indicating that irrelevant features become more problematic in low-data regimes.

4 Conclusion

We introduced the central distribution, a minimax-KL representative of a set, or more generally a hierarchy, of probability distributions. This provides a principled mechanism for converting partial or set-valued probabilistic knowledge into a single probability distribution suitable for Bayesian inference.

We showed that central distributions admit optimization-theoretic characterizations through convex-concave saddle point problems and established sufficient conditions for centrality based on the equalization of KL divergences. We further introduced central divergence as a measure of representational quality and derived bounds that enable the approximation of central distributions.

We further extended the framework to parametric families through the notion of a central prior and identified a sufficient fixed-point condition connecting central priors to Bernardo’s reference-prior methodology. We also proposed a notion of asymptotic

centrality for infinite sequences of random variables, providing a link between finite-sample and asymptotic constructions.

We showed that uncertainty over model structure, such as uncertainty about relevant feature subsets, can be represented through hierarchical collections of distributions. Central distributions then provide a principled mechanism for aggregating such structural uncertainty into a single predictive model.

Several important questions remain open. In particular, existence and uniqueness results for general hierarchies, computational methods for large-scale problems, and a rigorous asymptotic theory for central priors require further investigation. We hope that the central distribution framework provides a useful foundation for studying objective Bayesian inference, hierarchical uncertainty, and feature selection from a unified information-theoretic perspective.

References

- [1] Bernardo, J. M. (1979). “Reference Posterior Distributions for Bayesian Inference.” *Journal of the Royal Statistical Society: Series B (Methodological)*, 41(2): 113–128.
URL <https://doi.org/10.1111/j.2517-6161.1979.tb01066.x> 3
- [2] — (2005). “Reference Analysis.” In Dey, D. and Rao, C. (eds.), *Bayesian Thinking*, volume 25 of *Handbook of Statistics*, 17–90. Elsevier.
URL <https://www.sciencedirect.com/science/article/pii/S0169716105250022> 2, 3, 16, 17, 19
- [3] Boyd, S. and Vandenberghe, L. (2004). *Convex Optimization*. Cambridge University Press.
URL <https://books.google.com/books?id=IUZdAAAAQBAJ> 10
- [4] Jaynes, E. T. (1957). “Information Theory and Statistical Mechanics.” *Phys. Rev.*, 106: 620–630.
URL <https://link.aps.org/doi/10.1103/PhysRev.106.620> 3
- [5] Jeffreys, H. (1946). “An invariant form for the prior probability in estimation problems.” *Proceedings of the Royal Society of London. A. Mathematical and Physical Sciences*, 186(1007): 453–461.
URL <https://doi.org/10.1098/rspa.1946.0056> 3
- [6] Kullback, S. and Leibler, R. A. (1951). “On Information and Sufficiency.” *The Annals of Mathematical Statistics*, 22(1): 79 – 86.
URL <https://doi.org/10.1214/aoms/1177729694> 3
- [7] Rios Insua, D. and Ruggeri, F. (2000). *Robust Bayesian Analysis*, volume 152 of *Lecture Notes in Statistics*. New York: Springer-Verlag. 3
- [8] Walley, P. (1991). *Statistical Reasoning with Imprecise Probabilities*. Chapman & Hall/CRC Monographs on Statistics & Applied Probability. Taylor & Francis.
URL <https://books.google.com/books?id=-hbvAAAAAAAJ> 3